

# Chapter 1

## $n$ -Armed bandits

### 1.1 Notes

#### 1.1.1 $n$ -Armed Bandit Problem

We have  $n$  different options (actions) representing  $n$  different slot machines. Each action has a given reward, sampled from a stationary probability  $q(a)$  only dependent on the chosen action  $a$ . We want to maximize the (expected) total reward over a given (large) time  $T$ :  $\sum_{t=1}^T R_t$ . To do that, we estimate the value  $Q_t(a)$  of each action given what we have seen so far. Let  $R_t$  the reward at time  $t$  and  $N_t(a)$  the number of times the action  $a$  has been chosen so far.

#### 1.1.2 Estimating value

We estimate the value with:

$$Q_t(a) = \frac{R_1 + \dots + R_{N_t(a)}}{N_t(a)}$$

with  $Q_t(a) = Q_1(a)$  a default value. With  $N_t(a) \rightarrow \infty$  we have  $Q_t(a) \rightarrow q(a)$ .

Step-by-step, this can be calculated using incremental implementation to save computation time:

$$\begin{aligned} Q_{k+1} &= \frac{1}{k} \sum_{i=1}^k R_i \\ &\dots \\ Q_{k+1} &= Q_k + \frac{1}{k} (R_k - Q_k) \end{aligned}$$

This looks like  $NewEstimate \leftarrow OldEstimate + StepSize (Target - OldEstimate)$ , with  $StepSize = \frac{1}{k}$  here.

For tasks that never stop this estimation diverges, plus we may be interested in tracking a nonstationary problem. To achieve this, we can introduce a constant step size, that effectively weights recent rewards more heavily:

$$\begin{aligned} Q_{k+1} &= Q_k + \alpha (R_k - Q_k) \\ &\dots \\ Q_{k+1} &= (1 - \alpha)^k Q_1 + \alpha \sum_{i=1}^k (1 - \alpha)^{k-i} R_i \end{aligned}$$

As it turns out, this defines a weighted average with weights  $(1 - \alpha)^k, \alpha(1 - \alpha)^{k-1}, \dots, \alpha(1 - \alpha)^0$  (they sum to 1).

By denoting  $\alpha_k(a)$  the weight (step-size) used for the  $k$ -th selection of action  $a$ , we need to have two conditions:

1.  $\sum_{k=1}^{\infty} \alpha_k(a) = \infty$ , to guarantee that we overcome initial estimate, and
2.  $\sum_{k=1}^{\infty} \alpha_k^2(a) < \infty$ , to guarantee convergence.

## Choosing actions

To choose the action, the *greedy* way is to select the one with the highest value:  $A_t = \operatorname{argmax}_a Q_t(a)$ . Problem: this does not spend any time to sample other actions to refine the estimates  $Q_t(a)$ .

First solution:  $\epsilon$ -greedy algorithms, where  $A_t = \operatorname{argmax}_a Q_t(a)$   $1 - \epsilon$  of the times and  $A_t = \operatorname{uniform}(a)$  the other  $\epsilon$  of the times.

Second solution: optimistic initial values, to preferentially select unsampled actions.

Third solution: *Upper Confidence Bound (UCB)* action selection, with

$$A_t = \operatorname{argmax}_a \left( Q_t(a) + c \sqrt{\frac{\ln t}{N_t(a)}} \right)$$

Rationale: the square-root term is a measure of the uncertainty (or variance) in the estimate of the value of  $a$ . The higher this term (and  $c$ ), the higher the chance the action will be taken instead of the optimal action. UCB is often hard to transpose outside of the  $n$ -armed bandit problem.

## Gradient bandits

Instead of estimating the value of each action, we can estimate the relative preference  $H_t(a)$  of one action over others. We then compute the probability of taking action  $a$  using the softmax distribution (softmax "normalizes" its inputs so that any constant added to all preferences has no effect on the probability):

$$Pr(A_t = a) = \frac{e^{H_t(a)}}{\sum_{b=1}^n e^{H_t(b)}} = \pi_t(a)$$

Initially,  $H_1(a) = 0$ , so every action has the same probability to be chosen. We then use a variation on stochastic gradient ascent to update the preference after each step ( $A_t$  is the action taken at step  $t$  and  $R_t$  the corresponding reward):

$$\begin{aligned} H_{t+1}(A_t) &= H_t(A_t) + \alpha(R_t - \bar{R}_t)(1 - \pi_t(A_t)) && \text{and} \\ H_{t+1}(a) &= H_t(a) + \alpha(R_t - \bar{R}_t)\pi_t(A_t) && \forall a \neq A_t \end{aligned}$$

Explanation: if the reward ( $R_t$ ) is better than the average reward observed thus far ( $\bar{R}_t$ ), we increase the probability to select  $A_t$  and decrease the probabilities of all other actions in proportion to that difference ( $R_t - \bar{R}_t$ ). On the contrary, if the reward is lower than the current average reward, we decrease  $\pi_t(A_t)$  and increase all others. We increase/decrease more if the selected action had lower probability (thus the  $1 - \pi_t(A_t)$  term).

## 1.2 Exercises

**Exercise 2.1** *In the comparison in fig 2.1, which method will perform best in the long run in terms of cumulative reward, and cumulative probability of selecting the best action?*

Because of the law of large numbers, in the long run we have  $Q_t(a) \approx q(a)$  for  $\epsilon$ -greedy methods, since each actions will have been sampled a very large number of times. This is of course not true for the greedy method which correctly sample only one action, the one it always choses. For  $\epsilon$ -greedy method, in the long run we select the action  $A^* = \operatorname{argmax}_a q(a)$   $(1 - \epsilon)$  of the time and a random action  $\epsilon$  of the time. There doesn't seem to be any exact function that gives the expectation of the maximum of  $n$  iid normal variables, I could only find inequalities for large  $n$ ... I computed the value for  $n = 10$  using 1 million series of 10 iid normal variables, and got a value of  $\approx 1.54$ . The average reward when selecting action  $A^*$  is  $\mathbb{E}[q(A^*)] \approx 1.54$ . Meanwhile,

let  $Z$  be the reward when selection a random action, and  $Y \sim \mathcal{N}(0, 1)$  the noise term added to  $q(a)$  when computing the reward. We have:

$$\begin{aligned}\mathbb{E}[Z] &= \mathbb{E}[\mathbb{E}[q(a)] + Y] \\ &= \mathbb{E}[q(a)] + \mathbb{E}[Y] \\ &= 0\end{aligned}$$

since both random variables  $q(a)$  and  $Y$  follow a standard normal distribution. All in all, for  $\epsilon$ -greedy method, the average expected reward is:

$$\begin{aligned}\mathbb{E}[\bar{R}] &= \epsilon \mathbb{E}[Z] + (1 - \epsilon) \mathbb{E}[q(A^*)] \\ &\approx (1 - \epsilon) \times 1.54\end{aligned}$$

Which means 1.39 for  $\epsilon = 0.1$  and 1.52 for  $\epsilon = 0.01$ .

For the true greedy method, the first action  $A^\dagger$  for which the reward gets over 0 gets chosen everytime (at first approximation). Thus we need to find  $\mathbb{E}[q(A^\dagger)] = \mathbb{E}[q(a) | q(a) + Y > 0]$ . I could not manage to find a seemingly correct formula unfortunately (but this one should not be too hard for a statistician), so I once again computed an estimate for this number for  $n = 10$ , and came up with  $\mathbb{E}[q(A^\dagger)] \approx 0.56$ . This looks a bit weird since in fig 2.1 the average reward of the greedy method is around 1.

## Exercise 2.2