

Chapter 1

n -Armed bandits

1.1 Notes

1.1.1 n -Armed Bandit Problem

We have n different options (actions) representing n different slot machines. Each action has a given reward, sampled from a stationary probability $q(a)$ only dependent on the chosen action a . We want to maximize the (expected) total reward over a given (large) time T : $\sum_{t=1}^T R_t$. To do that, we estimate the value $Q_t(a)$ of each action given what we have seen so far. Let R_t the reward at time t and $N_t(a)$ the number of times the action a has been chosen so far.

1.1.2 Estimating value

We estimate the value with:

$$Q_t(a) = \frac{R_1 + \dots + R_{N_t(a)}}{N_t(a)}$$

with $Q_t(a) = Q_1(a)$ a default value. With $N_t(a) \rightarrow \infty$ we have $Q_t(a) \rightarrow q(a)$.

Step-by-step, this can be calculated using incremental implementation to save computation time:

$$\begin{aligned} Q_{k+1} &= \frac{1}{k} \sum_{i=1}^k R_i \\ &\dots \\ Q_{k+1} &= Q_k + \frac{1}{k} (R_k - Q_k) \end{aligned}$$

This looks like $NewEstimate \leftarrow OldEstimate + StepSize (Target - OldEstimate)$, with $StepSize = \frac{1}{k}$ here.

For tasks that never stop this estimation diverges, plus we may be interested in tracking a nonstationary problem. To achieve this, we can introduce a constant step size, that effectively weights recent rewards more heavily:

$$\begin{aligned} Q_{k+1} &= Q_k + \alpha (R_k - Q_k) \\ &\dots \\ Q_{k+1} &= (1 - \alpha)^k Q_1 + \alpha \sum_{i=1}^k (1 - \alpha)^{k-i} R_i \end{aligned}$$

As it turns out, this defines a weighted average with weights $(1 - \alpha)^k, \alpha(1 - \alpha)^{k-1}, \dots, \alpha(1 - \alpha)^0$ (they sum to 1).

By denoting $\alpha_k(a)$ the weight (step-size) used for the k -th selection of action a , we need to have two conditions:

1. $\sum_{k=1}^{\infty} \alpha_k(a) = \infty$, to guarantee that we overcome initial estimate, and
2. $\sum_{k=1}^{\infty} \alpha_k^2(a) < \infty$, to guarantee convergence.

Choosing actions

To choose the action, the *greedy* way is to select the one with the highest value: $A_t = \operatorname{argmax}_a Q_t(a)$. Problem: this does not spend any time to sample other actions to refine the estimates $Q_t(a)$.

First solution: ϵ -greedy algorithms, where $A_t = \operatorname{argmax}_a Q_t(a)$ $1 - \epsilon$ of the times and $A_t = \operatorname{uniform}(a)$ the other ϵ of the times.

Second solution: optimistic initial values, to preferentially select unsampled actions.

Third solution: *Upper Confidence Bound (UCB)* action selection, with

$$A_t = \operatorname{argmax}_a \left(Q_t(a) + c \sqrt{\frac{\ln t}{N_t(a)}} \right)$$

Rationale: the square-root term is a measure of the uncertainty (or variance) in the estimate of the value of a . The higher this term (and c), the higher the chance the action will be taken instead of the optimal action. UCB is often hard to transpose outside of the n -armed bandit problem.

Gradient bandits

Instead of estimating the value of each action, we can estimate the relative preference $H_t(a)$ of one action over others. We then compute the probability of taking action a using the softmax distribution (softmax "normalizes" its inputs so that any constant added to all preferences has no effect on the probability):

$$Pr(A_t = a) = \frac{e^{H_t(a)}}{\sum_{b=1}^n e^{H_t(b)}} = \pi_t(a)$$

Initially, $H_1(a) = 0$, so every action has the same probability to be chosen. We then use a variation on stochastic gradient ascent to update the preference after each step (A_t is the action taken at step t and R_t the corresponding reward):

$$\begin{aligned} H_{t+1}(A_t) &= H_t(A_t) + \alpha(R_t - \bar{R}_t)(1 - \pi_t(A_t)) && \text{and} \\ H_{t+1}(a) &= H_t(a) + \alpha(R_t - \bar{R}_t)\pi_t(A_t) && \forall a \neq A_t \end{aligned}$$

Explanation: if the reward (R_t) is better than the average reward observed thus far (\bar{R}_t), we increase the probability to select A_t and decrease the probabilities of all other actions in proportion to that difference ($R_t - \bar{R}_t$). On the contrary, if the reward is lower than the current average reward, we decrease $\pi_t(A_t)$ and increase all others. We increase/decrease more if the selected action had lower probability (thus the $1 - \pi_t(A_t)$ term).

1.2 Exercises

Exercise 2.1 *In the comparison in fig 2.1, which method will perform best in the long run in terms of cumulative reward, and cumulative probability of selecting the best action?*

Because of the law of large numbers, in the long run we have $Q_t(a) \approx q(a)$ for ϵ -greedy methods, since each actions will have been sampled a very large number of times. This is of course not true for the greedy method which correctly sample only one action, the one it always choses. For ϵ -greedy method, in the long run we select the action $A^* = \operatorname{argmax}_a q(a)$ $(1 - \epsilon)$ of the time and a random action ϵ of the time. There doesn't seem to be any exact function that gives the expectation of the maximum of n iid normal variables, I could only find inequalities for large n ... I computed the value for $n = 10$ using 1 million series of 10 iid normal variables, and got a value of ≈ 1.54 . The average reward when selecting action A^* is $\mathbb{E}[q(A^*)] \approx 1.54$. Meanwhile,

let Z be the reward when selection a random action, and $Y \sim \mathcal{N}(0, 1)$ the noise term added to $q(a)$ when computing the reward. We have:

$$\begin{aligned}\mathbb{E}[Z] &= \mathbb{E}[\mathbb{E}[q(a)] + Y] \\ &= \mathbb{E}[q(a)] + \mathbb{E}[Y] \\ &= 0\end{aligned}$$

since both random variables $q(a)$ and Y follow a standard normal distribution. All in all, for ϵ -greedy method, the average expected reward is:

$$\begin{aligned}\mathbb{E}[\bar{R}] &= \epsilon \mathbb{E}[Z] + (1 - \epsilon) \mathbb{E}[q(A^*)] \\ &\approx (1 - \epsilon) \times 1.54\end{aligned}$$

Which means 1.39 for $\epsilon = 0.1$ and 1.52 for $\epsilon = 0.01$.

For the true greedy method, the first action A^\dagger for which the reward gets over 0 gets chosen everytime (at first approximation). Thus we need to find $\mathbb{E}[q(A^\dagger)] = \mathbb{E}[q(a) | q(a) + Y > 0]$. I could not manage to find a seemingly correct formula unfortunately (but this one should not be too hard for a statistician), so I once again computed an estimate for this number for $n = 10$, and came up with $\mathbb{E}[q(A^\dagger)] \approx 0.56$. This looks a bit weird since in fig 2.1 the average reward of the greedy method is around 1.

For ϵ -greedy method, the probability of selecting the best action when we decided to chose a random action is obviously $\frac{1}{n}$. As for when we decide to chose the action with the highest estimated value, I get stuck. We need to find, given $(Y, Z) \sim \mathcal{N}(0, 1)^2$, the probability $Pr\{q(a) + Y > q(A^*) + Z\}$. That honestly seems daunting to me and I would not even know where to begin anyway... Graphically, the curve seems to plateau at around 80% for $\epsilon = 0.1$, and should be plateauing over this value for $\epsilon = 0.01$. Same for the greedy method, which seems to plateau at around 35%.

Exercise 2.2 Give pseudocode for a complete algorithm for the n -armed bandit problem. Use greedy action selection and incremental computation of action values with $\alpha = \frac{1}{k}$ step-size parameter. Assume a function `bandit(a)` that takes an action and returns a reward. Use arrays and variables; do not subscript anything by the time index t . Indicate how the action values are initialized and updates after each reward. Indicate how the step-size parameters are set for each action as a function of how many times it has been tried.

1. Initialization
 - $Q(a) \leftarrow 0$ for all $a \in \mathcal{A}$
 - $k(a) \leftarrow 1$ for all $a \in \mathcal{A}$

2. Action selection and value update

Repeat

$A \leftarrow \operatorname{argmax}_a Q(a)$ (resolve ties randomly)

$R \leftarrow \text{bandit}(A)$

$Q(A) \leftarrow Q(A) + \frac{1}{k}(R - Q(A))$

$k(A) \leftarrow k(A) + 1$

until end of epoch

Exercise 2.3 *If the step-size parameters, α_k , are not constant, then the estimate Q_k is a weighted average of previously received rewards with a weighting different from that given by (2.6). What is the weighting on each prior reward for the general case, analogous to (2.6), in terms of α_k ?*

$$\begin{aligned} Q_{k+1} &= Q_k + \alpha_k(R_k - Q_k) \\ &= \alpha_k R_k + (1 - \alpha_k)Q_k \\ &= \alpha_k R_k + (1 - \alpha_k)[\alpha_{k-1}R_{k-1} + (1 - \alpha_{k-1})Q_{k-1}] \\ &= \alpha_k R_k + \alpha_{k-1}(1 - \alpha_k)R_{k-1} + (1 - \alpha_k)(1 - \alpha_{k-1})Q_{k-1} \\ &= Q_1 \prod_{i=1}^k (1 - \alpha_i) + \sum_{i=1}^k \left(\alpha_i R_i \prod_{j=i+1}^k (1 - \alpha_j) \right) \end{aligned}$$

The weights are then $\prod_{i=1}^k (1 - \alpha_i)$ for Q_1 and $\alpha_i R_i \prod_{j=i+1}^k (1 - \alpha_j)$ for R_i .

Exercise 2.4 (programming experiment) *Design and conduct an experiment to demonstrate the difficulties that sample-average methods have for nonstationary problems. Use a modified version of the 10-armed testbed in which all the $q(a)$ start out equal and then take independent random walks. Prepare plots like Figure 2.1 for an action-value method using sample averages, incrementally computed by $\alpha = \frac{1}{k}$, and another action-value method using a constant step-size parameter, $\alpha = 0.1$. Use $\epsilon = 0.1$ and, if necessary, runs longer than 1000 plays.*

See file `exercise_2_4.py` for the Python code. The code is composed with 3 classes (2 for bandit problem definition, `Bandit` and `DriftBandit`, and 1 for a general action-value method, `EpsGreedyLearner`) and 2 functions (`run_drift` and `experiments`) to conduct the experiment.

Figure 1.1.left shows that a fixed step-size consistently outperforms the "average" step-size. Figure 1.1.right clearly shows the difficulty of the action-value method to track a non-stationary problem. Indeed, the Figure 2.1 in




Figure 1.1: Left: average rewards (blue: $\alpha = \frac{1}{k}$, green: $\alpha = 0.1$, red: green/blue). Right: probability of selecting the best action (blue: $\alpha = \frac{1}{k}$, green: $\alpha = 0.1$).

the original manuscript reports probabilities of selecting the best action to be around 80% in the long-run for $\epsilon = 0.1$. For the same ϵ value, this method only reports probabilities of 50-60% – although this number could still increase with a higher number of time steps.